

flight-price-prediction

April 3, 2024

1 Flight Price Prediction

In this notebook, we will consider the problem of modelling flight price prediction based on the data from Kaggle website.

1.1 Import Necessary Libraries

First, we need to import the libraries that will be used throughout this notebook.

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn as skl
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.feature_selection import mutual_info_regression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_predict
from sklearn.model_selection import cross_validate
```

1.2 Load Datasets

Now, let's load the datasets that we will be using for our analysis.

```
[ ]: # Load datasets
business_df = pd.read_csv('../datasets/business.csv')
economy_df = pd.read_csv('../datasets/economy.csv')
clean_dataset = pd.read_csv('../datasets/Clean_Dataset.csv')
business_df.head()
economy_df.head()
clean_dataset.head()
```

```
[ ]: Unnamed: 0  airline  flight source_city departure_time stops \
0          0  SpiceJet  SG-8709      Delhi      Evening  zero
1          1  SpiceJet  SG-8157      Delhi  Early_Morning  zero
2          2  AirAsia   I5-764      Delhi  Early_Morning  zero
3          3  Vistara   UK-995      Delhi      Morning  zero
4          4  Vistara   UK-963      Delhi      Morning  zero
```

```
      arrival_time destination_city  class  duration  days_left  price
0          Night      Mumbai  Economy    2.17         1  5953
1          Morning      Mumbai  Economy    2.33         1  5953
2  Early_Morning      Mumbai  Economy    2.17         1  5956
3      Afternoon      Mumbai  Economy    2.25         1  5955
4          Morning      Mumbai  Economy    2.33         1  5955
```

```
[ ]: clean_dataset.shape
      clean_dataset.describe(include='all')
```

```
[ ]: Unnamed: 0  airline  flight source_city departure_time  stops \
count  300153.000000  300153  300153      300153      300153  300153
unique          NaN         6    1561         6         6         3
top          NaN  Vistara  UK-706      Delhi      Morning      one
freq          NaN  127859    3235      61343      71146  250863
mean   150076.000000         NaN      NaN      NaN      NaN      NaN
std    86646.852011         NaN      NaN      NaN      NaN      NaN
min         0.000000         NaN      NaN      NaN      NaN      NaN
25%    75038.000000         NaN      NaN      NaN      NaN      NaN
50%   150076.000000         NaN      NaN      NaN      NaN      NaN
75%   225114.000000         NaN      NaN      NaN      NaN      NaN
max   300152.000000         NaN      NaN      NaN      NaN      NaN
```

```
      arrival_time destination_city  class  duration  days_left \
count      300153      300153  300153  300153.000000  300153.000000
unique         6         6         2         NaN         NaN
top      Night      Mumbai  Economy         NaN         NaN
freq      91538      59097  206666         NaN         NaN
mean         NaN         NaN      NaN    12.221021    26.004751
std         NaN         NaN      NaN     7.191997    13.561004
min         NaN         NaN      NaN     0.830000     1.000000
25%         NaN         NaN      NaN     6.830000    15.000000
50%         NaN         NaN      NaN    11.250000    26.000000
75%         NaN         NaN      NaN    16.170000    38.000000
max         NaN         NaN      NaN    49.830000    49.000000
```

```
      price
count  300153.000000
unique      NaN
top      NaN
```

```

freq          NaN
mean    20889.660523
std     22697.767366
min      1105.000000
25%      4783.000000
50%      7425.000000
75%     42521.000000
max     123071.000000

```

```
[ ]: clean_dataset.dropna(inplace=True)
clean_dataset.shape
```

```
[ ]: (300153, 12)
```

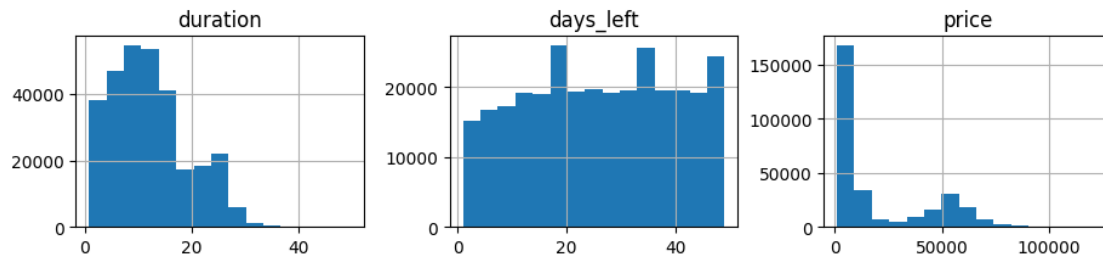
```
[ ]: clean_dataset.isnull().sum()
```

```
[ ]: Unnamed: 0      0
airline            0
flight            0
source_city       0
departure_time    0
stops            0
arrival_time      0
destination_city  0
class             0
duration          0
days_left        0
price            0
dtype: int64
```

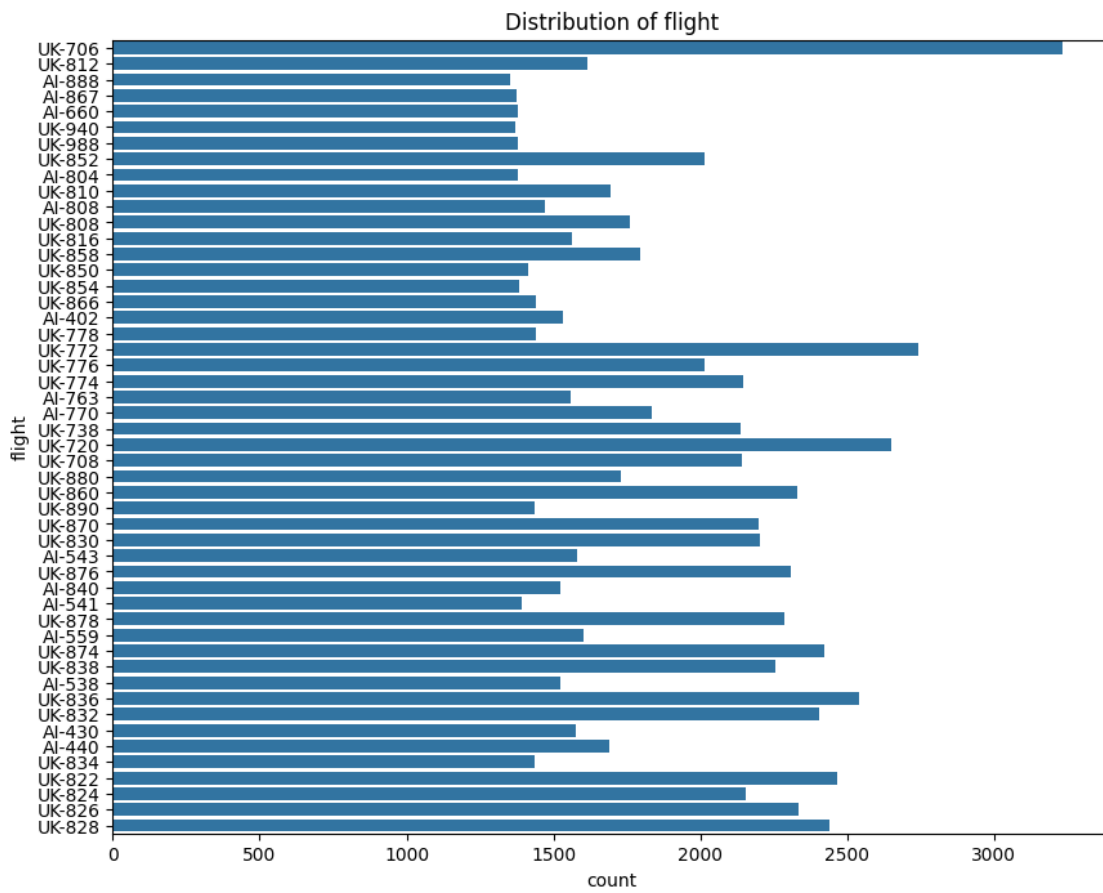
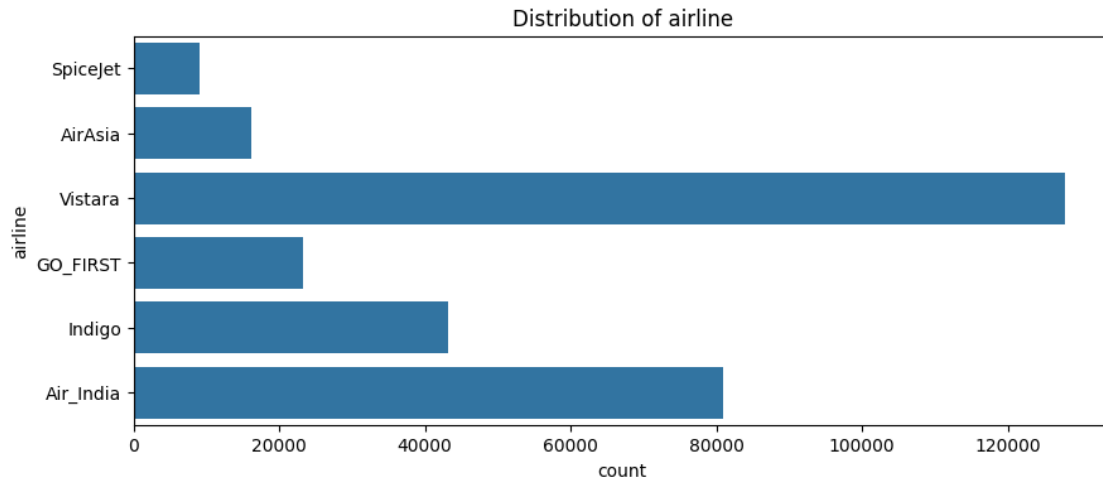
2 Let's visualize the first few rows of the dataset

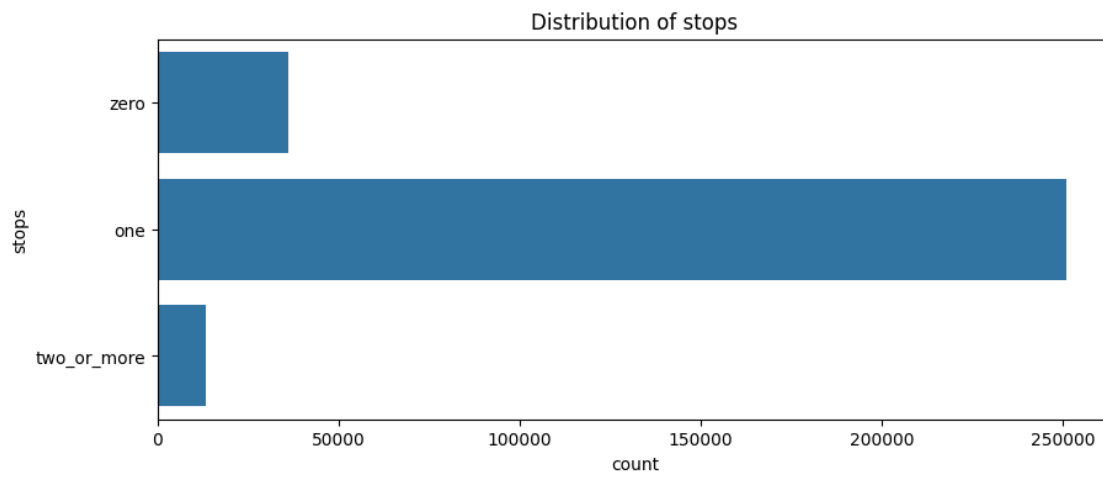
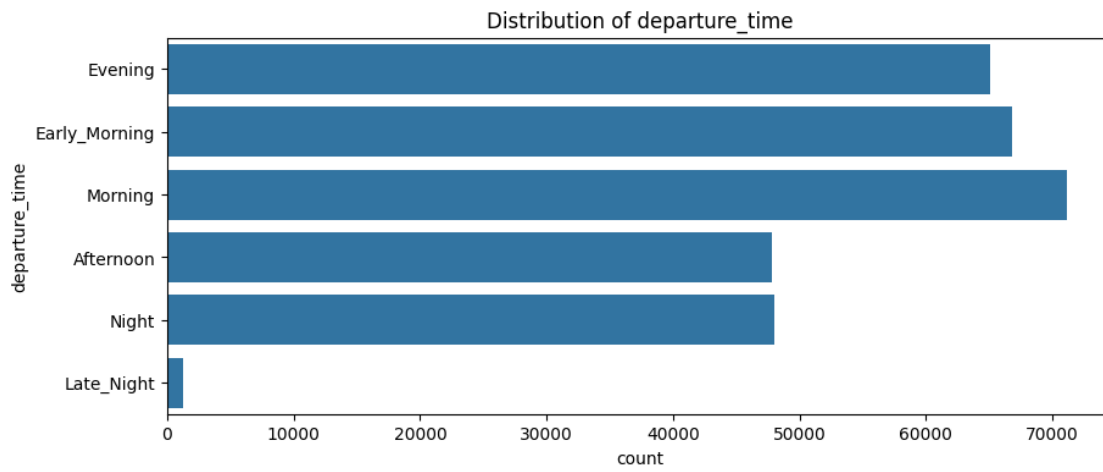
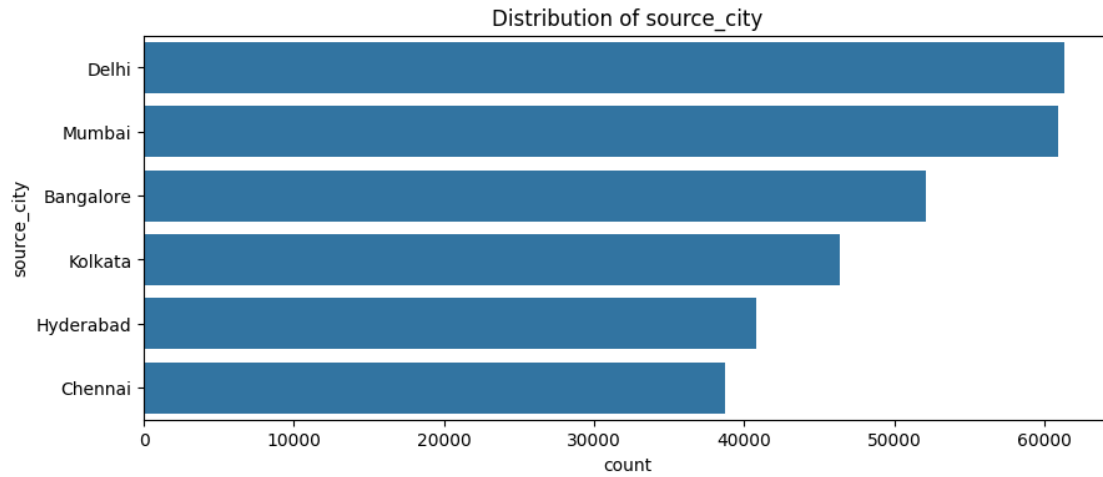
```
[ ]: # Plotting histograms for all numeric features to understand distributions
# exclude the unnamed column
clean_dataset.drop('Unnamed: 0', axis=1, inplace=True)
clean_dataset.hist(bins=15, figsize=(15, 10), layout=(4, 4))
plt.suptitle('Histograms of numeric features')
plt.show()
```

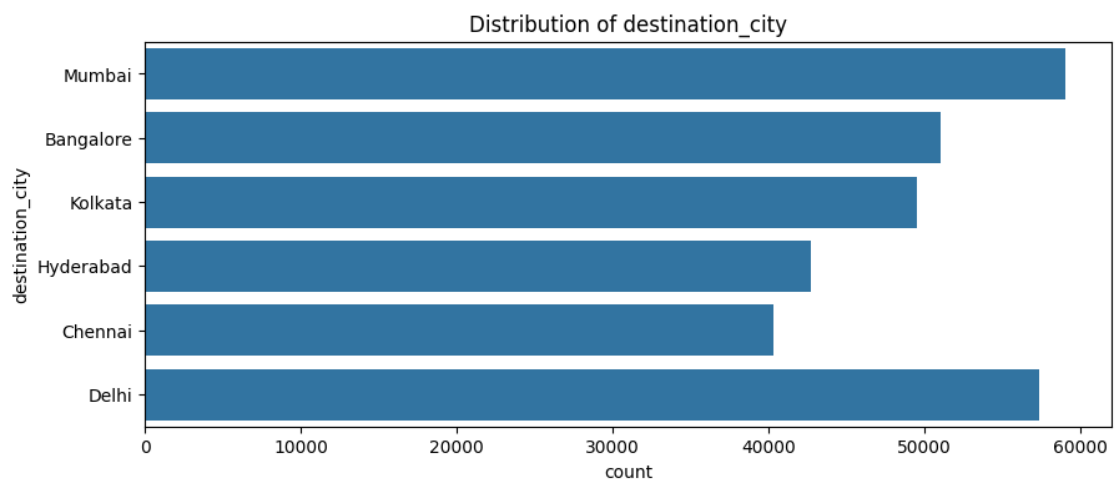
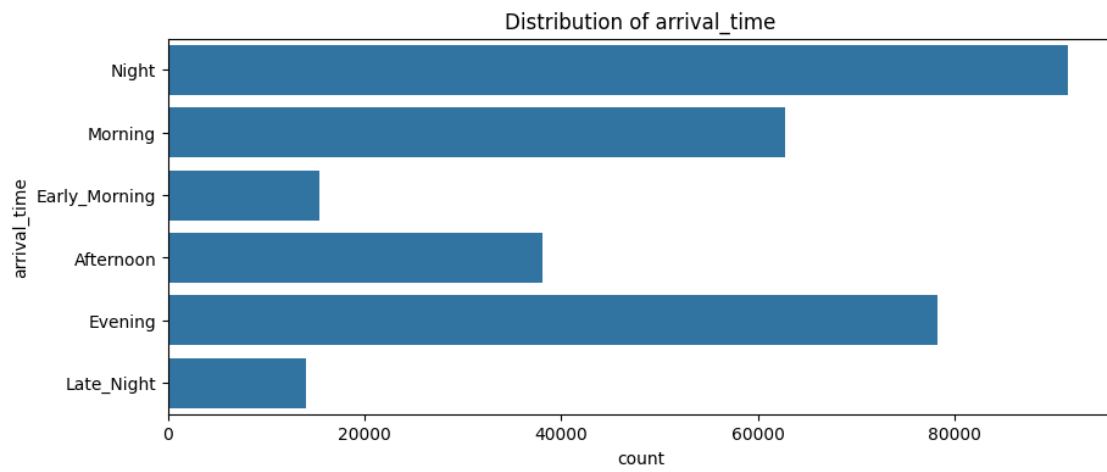
Histograms of numeric features

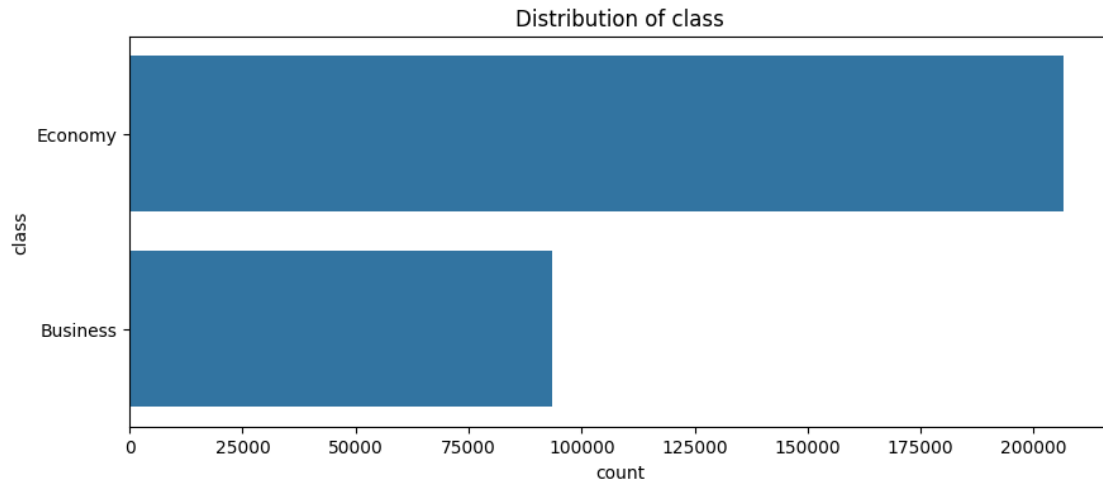


```
[ ]: # For categorical data, we can use count plots to understand the distribution
      ↪ of categories
for column in clean_dataset.select_dtypes(include=['object']).columns:
    # Plotting count plots for all categorical features
    # If the number of categories is too high, e.g., flight, we can filter the
    ↪ top 50 categories to make the plot more readable
    if column != 'flight':
        plt.figure(figsize=(10, 4))
        sns.countplot(y=column, data=clean_dataset)
        plt.title(f'Distribution of {column}')
        plt.show()
    else:
        top_categories = clean_dataset[column].value_counts().index[:50] # Get
        ↪ top 50 categories
        filtered_data = clean_dataset[clean_dataset[column].
        ↪ isin(top_categories)]
        plt.figure(figsize=(10, 8))
        sns.countplot(y=column, data=filtered_data)
        plt.yticks(fontsize=10)
        plt.title(f'Distribution of {column}')
        plt.show()
```

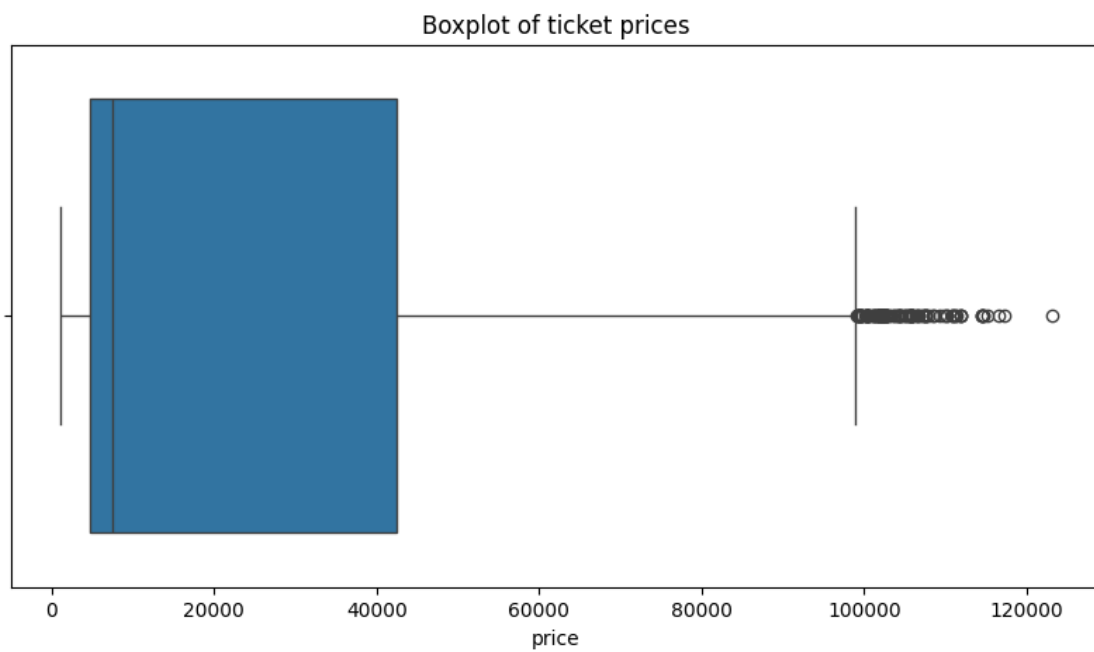








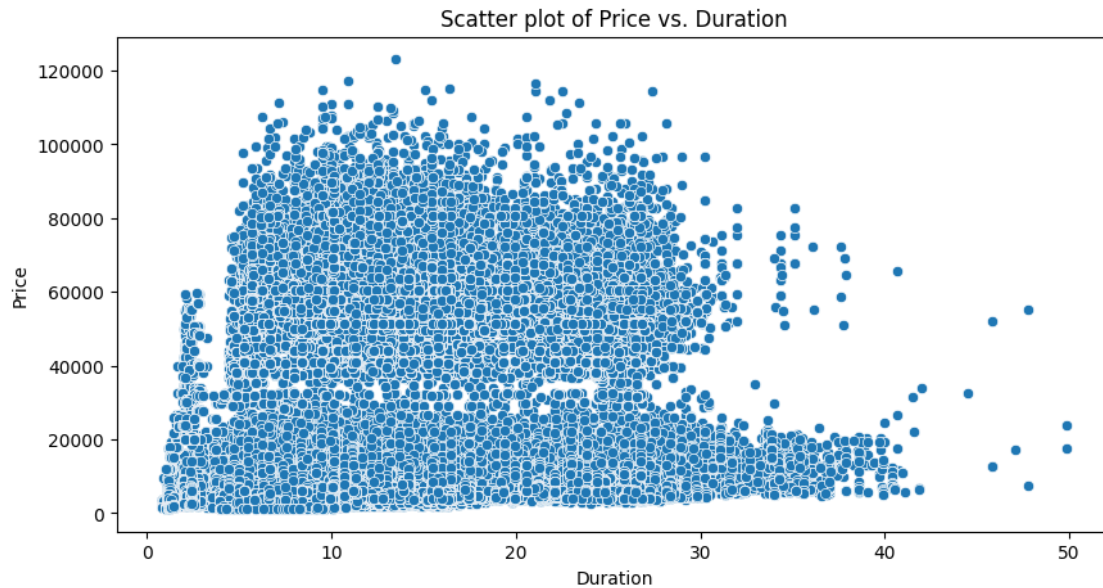
```
[ ]: # Boxplot for the price column to see its distribution and spot any outliers
plt.figure(figsize=(10, 5))
sns.boxplot(x=clean_dataset['price'])
plt.title('Boxplot of ticket prices')
plt.show()
```



```
[ ]: # A scatter plot to visualize the relationship between two variables, for
      ↪ example, price and duration
```

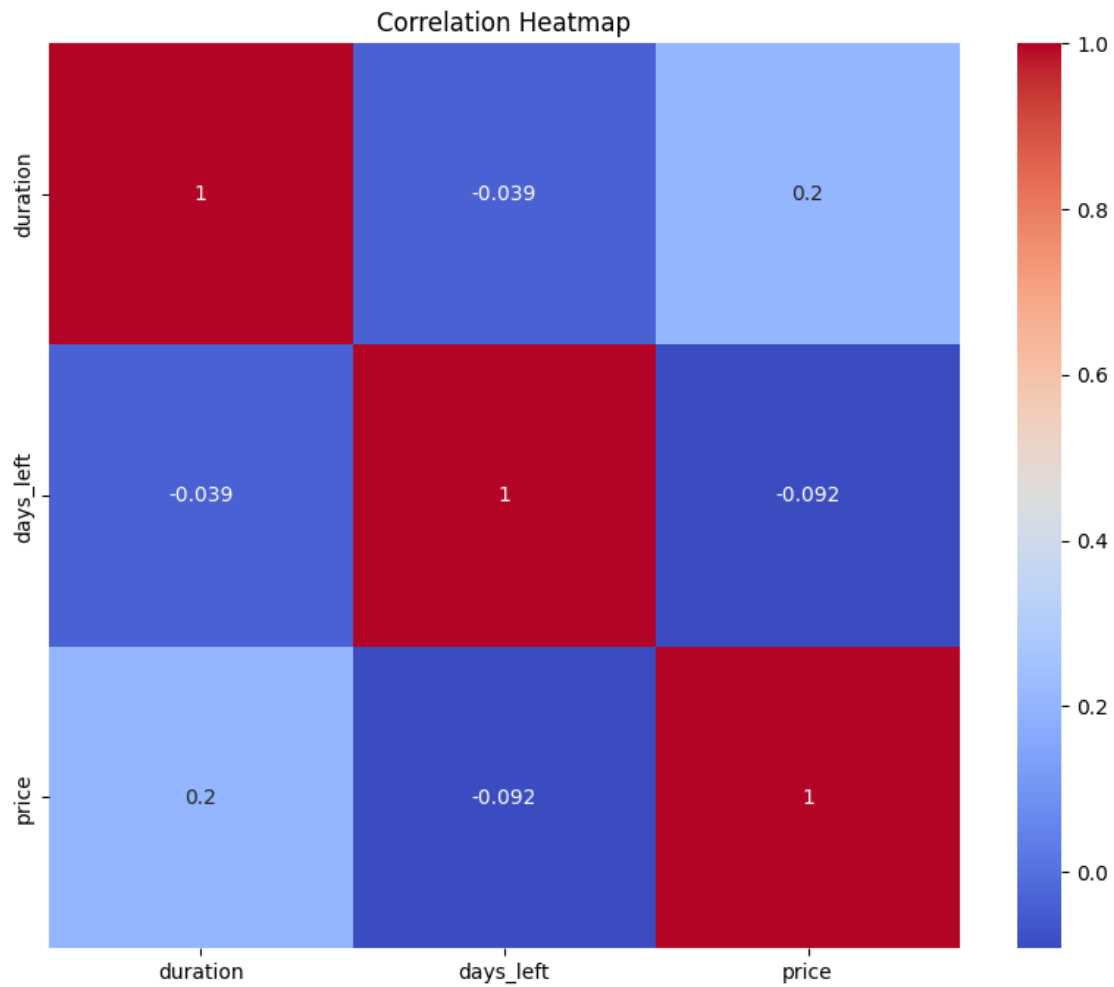


```
plt.figure(figsize=(10, 5))
sns.scatterplot(x=clean_dataset['duration'], y=clean_dataset['price'])
plt.title('Scatter plot of Price vs. Duration')
plt.xlabel('Duration')
plt.ylabel('Price')
plt.show()
```



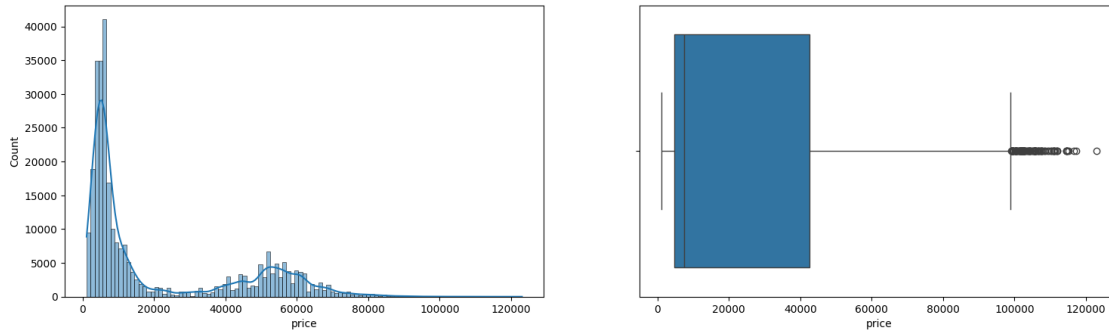
```
[ ]: # Correlation heatmap to understand the relationships between variables
# Select only the numeric columns for correlation
numeric_dataset = clean_dataset.select_dtypes(include=[np.number])
correlation_matrix = numeric_dataset.corr()
```

```
[ ]: # Visualize the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



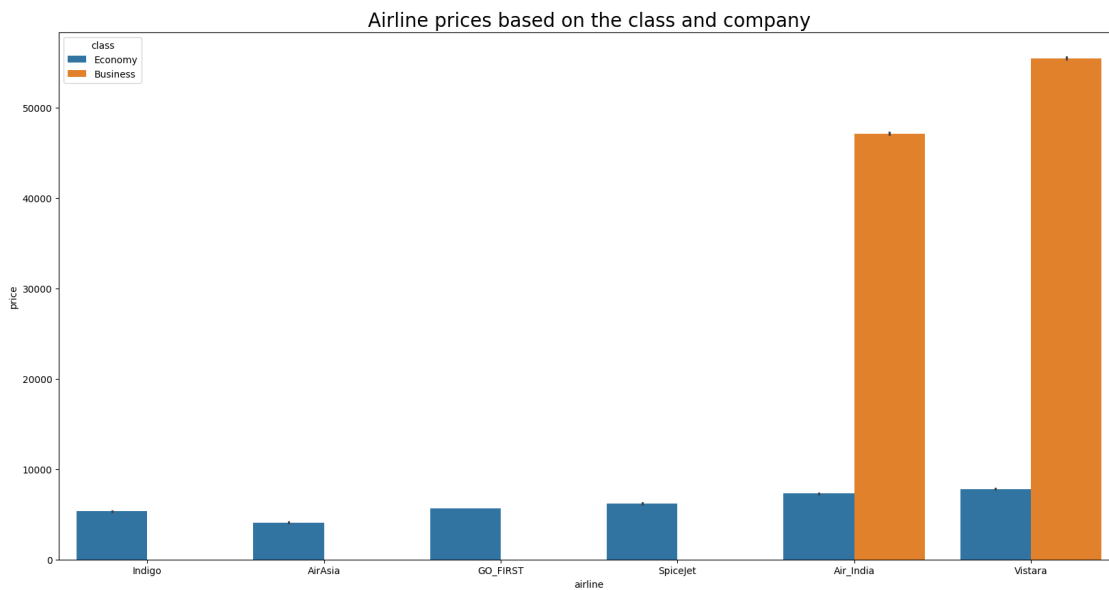
```
[ ]: plt.figure(figsize = (18,5))  
plt.subplot(1,2,1)  
sns.histplot(x = 'price', data = clean_dataset, kde = True)  
plt.subplot(1,2,2)  
sns.boxplot(x = 'price', data = clean_dataset)
```

```
[ ]: <Axes: xlabel='price'>
```



```
[ ]: plt.figure(figsize=(20, 10))
sns.barplot(x='airline',y='price',hue="class",data=clean_dataset.
↳sort_values("price")).set_title('Airline prices based on the class and_
↳company',fontsize=20)
```

```
[ ]: Text(0.5, 1.0, 'Airline prices based on the class and company')
```



3 Make data transformation

```
[ ]: transformed_dataset = clean_dataset.copy()
transformed_dataset['Economy'] = clean_dataset['class'] == 'Economy'
transformed_dataset.drop('class', axis=1, inplace=True)
```

```
[ ]: #transformed_dataset['source_city'].unique()
```

```
[ ]: city_size = { # this is for year 2011 - https://en.wikipedia.org/wiki/
    ↪List_of_cities_in_India_by_population
    'Delhi': 110,
    'Mumbai': 124,
    'Bangalore': 84,
    'Kolkata': 44,
    'Hyderabad': 69,
    'Chennai' : 46
}
transformed_dataset['source_size'] = transformed_dataset['source_city'].
    ↪replace(city_size)
transformed_dataset.drop('source_city', axis=1, inplace=True)
transformed_dataset['destination_size'] =
    ↪transformed_dataset['destination_city'].replace(city_size)
transformed_dataset.drop('destination_city', axis=1, inplace=True)
```

```
[ ]: transformed_dataset = pd.
    ↪get_dummies(transformed_dataset,columns=['departure_time','arrival_time'])
```

```
[ ]: stops = {
    'zero': 0,
    'one': 1,
    'two_or_more': 2,
}
transformed_dataset['stops_num'] = transformed_dataset['stops'].replace(stops)
transformed_dataset.drop('stops', axis=1, inplace=True)
```

```
[ ]: transformed_dataset = pd.get_dummies(transformed_dataset,columns=['airline'])
```

```
[ ]: transformed_dataset['flight_num'] = pd.
    ↪factorize(transformed_dataset['flight'])[0]
transformed_dataset.drop('flight', axis=1, inplace=True)
```

```
[ ]: transformed_dataset.head()
```

```
[ ]:
duration  days_left  price  Economy  source_size  destination_size  \
0        2.17         1   5953     True         110             124
1        2.33         1   5953     True         110             124
2        2.17         1   5956     True         110             124
3        2.25         1   5955     True         110             124
4        2.33         1   5955     True         110             124

departure_time_Afternoon  departure_time_Early_Morning  \
0                      False                          False
1                      False                          True
2                      False                          True
3                      False                          False
```

4		False		False
---	--	-------	--	-------

	departure_time_Evening	departure_time_Late_Night	...	\
0	True	False	...	
1	False	False	...	
2	False	False	...	
3	False	False	...	
4	False	False	...	

	arrival_time_Morning	arrival_time_Night	stops_num	airline_AirAsia	\
0	False	True	0	False	
1	True	False	0	False	
2	False	False	0	True	
3	False	False	0	False	
4	True	False	0	False	

	airline_Air_India	airline_GO_FIRST	airline_Indigo	airline_SpiceJet	\
0	False	False	False	True	
1	False	False	False	True	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	airline_Vistara	flight_num
0	False	0
1	False	1
2	False	2
3	True	3
4	True	4

[5 rows x 26 columns]

```
[ ]: transformed_dataset.describe()
# output the transformed dataset to a new CSV file
transformed_dataset.to_csv('../datasets/Transformed_Dataset.csv', index=False)
```